

- Analysis: cuSPARSE, cuBLAS, and GraphBLAST
 - Semiring (Refrencing GraphBLAST)
 - 1. cuBLAS: Dense Linear Algebra Primitive
 - Usage
 - 2. cuSPARSE: Low-Level Sparse Primitives
 - Usage
 - 3. GraphBLAST: Semiring-Optimized Graph Analytics
 - Architectural Advantages
 - Vector
 - Matrix
 - Operation
 - Publications
 - 4. Critical Comparison

Analysis: cuSPARSE, cuBLAS, and GraphBLAST

For Sparse Graph Algorithm Development on GPUs

Semiring (Refrencing GraphBLAST)

A semiring is the computation on vertex and edge of the graph. In standard matrix multiplication the semiring used is the (+, x) arithmetic semiring. In GraphBLAS, when the semiring is applied to this operation, it represents the transformation that is required to transform the input vector into the output vector. What the (+, x) represent are:

- x: computation per edge, generates up to num_edges intermediate elements
- +: computation in the reduction of intermediates back down to a subset of vertices, up to num_verts elements

The most frequently used semirings (with their common usage in brackets) are:

- PlusMultipliesSemiring: arithmetic semiring (classical linear algebra)
- LogicalOrAndSemiring: Boolean semiring (graph connectivity)
- MinimumPlusSemiring: tropical min-plus semiring (shortest path)

- **MaximumMultipliesSemiring**: tropical max-times semiring (maximal independent set)

1. cuBLAS: Dense Linear Algebra Primitive

Usage

- Optimized for **dense matrix operations** (GEMM, GEMV) using NVIDIA tensor cores.

Use Case: AI/ML dense tensor operations, not graph analytics

2. cuSPARSE: Low-Level Sparse Primitives

Usage

- GPU-accelerated SpMV/SpMM (30-150× faster than CPU)

Limitations:

- No native semiring support

Use Case: Sparse linear algebra in CFD/seismic, not graph abstractions.

3. GraphBLAST: Semiring-Optimized Graph Analytics

Architectural Advantages

3.1 Native Semiring Support Implements GraphBLAS standard with algebraic graph algorithms:

3.2 GPU-Specific Optimizations

- **Direction optimization:** Auto-switches push/pull traversal
- **Merge-based load balancing:**

```
if (avg_edges < 4) launch_coalesced_kernel();  
else launch_workstealing_kernel();
```

- **Code Example (Concise BFS):**

```
GrB_vxm(levels, GrB_NULL, GrB_LOR_LAND_SEMIRING_INT32, A, levels, GrB_NULL);
```

- **3.3 Single Source Shortest Path**

```
#include "graphblas/graphblas.hpp"  
  
// Single-source shortest-path on adjacency matrix A from source s  
graphblas::Info ssspSimple( Vector<float>*      v,  
                             const Matrix<float>* A,  
                             Index               s,  
                             Descriptor*         desc ) {  
  
    // Get number of vertices  
    graphblas::Index A_nrows;  
    A->nrows(&A_nrows);  
  
    // Distance vector (v)  
    std::vector<graphblas::Index> indices(1, s);  
    std::vector<float> values(1, 0.f);  
    v->build(&indices, &values, 1, GrB_NULL);  
  
    // Buffer vector (w)  
    graphblas::Vector<float> w(A_nrows);  
  
    // Semiring zero vector (zero)  
    graphblas::Vector<float> zero(A_nrows);  
    zero.fill(std::numeric_limits<float>::max());  
  
    // Initialize loop variables  
    graphblas::Index iter = 1;  
    float succ_last = 0.f;  
    float succ = 1.f;  
  
    do {  
        succ_last = succ;
```

```

// v = v + v * A^T (do relaxation on distance vector v)
graphblas::vxm<float, float, float, float>(&w, GrB_NULL, GrB_NULL,
    graphblas::MinimumPlusSemiring<float>(), v, A, desc);
graphblas::eWiseAdd<float, float, float, float>(v, GrB_NULL, GrB_NULL,
    graphblas::MinimumPlusSemiring<float>(), v, &w, desc);

// w = v < FLT_MAX (get all reachable vertices)
graphblas::eWiseMult<float, float, float, float>(&w, GrB_NULL, GrB_NULL,
    graphblas::PlusLessSemiring<float>(), v, &zero, desc);

// succ = reduce(w) (do reduction on all reachable distances)
graphblas::reduce<float, float>(&succ, GrB_NULL,
    graphblas::PlusMonoid<float>(), &w, desc);
iter++;

// Loop until total reachable distance has converged
} while (succ_last != succ);

return GrB_SUCCESS;
}

```

The idea behind GraphBLAS is that four concepts can be used to implement many graph algorithms: vector, matrix, operation and semiring.

Vector

A vector is a subset of vertices of some graph.

Matrix

A matrix is the adjacency matrix of some graph.

Operation

An operation is the memory access pattern common to most graph algorithms (equivalent Ligra terminology is shown in brackets):

- **mxv**: matrix-vector multiply (EdgeMap)
- **vxm**: vector-matrix multiply (EdgeMap)
- **mxm**: matrix-matrix multiply (multi-frontier EdgeMap)
- **eWiseAdd**: elementwise addition (VertexMap)
- **eWiseMult**: elementwise multiplication (VertexMap)

Publications

1. Carl Yang, Aydın Buluç, and John D. Owens. **GraphBLAST: A High-Performance Linear Algebra-based Graph Framework on the GPU**. arXiv preprint arXiv:1908.01407 (2019). [[arXiv](#)]
2. Carl Yang, Aydın Buluç, and John D. Owens. **Design Principles for Sparse Matrix Multiplication on the GPU**. In *Proceedings of the 24th International European Conference on Parallel and Distributed Computing*, Euro-Par, pages 672-687, August 2018. Distinguished Paper and Best Artifact Award. [[DOI](#) | [http](#) | [slides](#)]
3. Carl Yang, Aydın Buluç, John D. Owens. **Implementing Push-Pull Efficiently in GraphBLAS**. In *Proceedings of the International Conference on Parallel Processing*, ICPP, pages 89:1-89:11, August 2018. [[DOI](#) | [http](#) | [slides](#)]
4. Carl Yang, Yangzihao Wang, and John D. Owens. **Fast Sparse Matrix and Sparse Vector Multiplication Algorithm on the GPU**. In **Graph Algorithms Building Blocks**, In *Graph Algorithm Building Blocks*, GABB, pages 841–847, May 2015. [[DOI](#) | [http](#) | [slides](#)]

4. Critical Comparison

Aspect	cuBLAS	cuSPARSE	GraphBLAST
Target	Dense matrices	Sparse matrices	Sparse graphs
Semiring Support	None	Manual emulation	Native (GraphBLAS)
BFS Lines of Code	150+	80	25
Load Balancing	N/A	Manual	Auto-optimized

While cuSPARSE remains valuable for numerical sparse linear algebra, GraphBLAST’s focus on graph algorithm’s indicates it to be a better choice for now.

Sources:^[^1] [cuSPARSE Documentation](#)^[^2] [NVIDIA cuSPARSE Overview](#)^[^3] [Matrix Multiplication Benchmark \(arXiv:2405.17322\)](#)^[^4] [NVIDIA cuBLAS Documentation](#)^[^5]

[GraphBLAS Forum Collection](#)^[^6] [GraphBLAST Performance Analysis](#)^[^7] [UC Davis](#)

[GraphBLAST Report](#)^[^9] [GraphBLAST ACM Paper Summary](#)^[^10]

[SuiteSparse:GraphBLAS Implementation](#)